

SUBNET MAGAZINE

REPORTING · Issue 016, The Q2 retrospective

Ninety days that tested the Bittensor thesis

Between 12 February and 12 May 2026, Bittensor produced the most significant decentralized-training milestone to date, lost the team that produced it, watched community miners rebuild three top-emission subnets without coordination, and ships a governance upgrade to mainnet on May 13. The Q2 retrospective.

SUBNET MAGAZINE · LIVE

2026-05-12 · filed by Subnet Magazine

Between 12 February and 12 May 2026, Bittensor produced the most significant decentralized-training milestone to date, lost the team that produced it, watched community miners rebuild three top-emission subnets without coordination, and ships a governance upgrade to mainnet on May 13. The Q2 retrospective.

2026-05-12 · filed by Subnet Magazine

Article Search

Home @subnetmagazine Relevant people

Explore @subnetmagazine Bittensor Subnet Review Notifications

What's happening Follow OBSESSION Movie Chat In Theaters TODAY Promoted by Obsession

Grok Trending in United States

Premium Bittensor Thesis Only on X · Trending #007FirstLightRTX

Bookmarks 1 101 Sports · Trending Michael Vick

Articles Bittensor in May 2026: Ninety Days Show more

Profile That Tested the Thesis Terms of Service | Privacy Policy

Accessibility | Ads info | More Between February 12 and May 12, 2026, Bittensor produced the most significant technical milestone in decentralized AI training to date, lost

the team that produced it under circumstances that wiped roughly \$900 Post million in market capitalization in a single afternoon, watched community miners rebuild three of its highest-emission subnets from open-source code without central operator coordination, and is shipping a protocol-level governance upgrade to mainnet on May 13.

The network's response to its first mature governance crisis is the most important data point in decentralized AI this year. It also exposed structural problems the ecosystem had been slow to acknowledge, and accelerated their resolution faster than the alternative path would have allowed.

Everything about how Bittensor should be read against the closed lab economy depends on what it actually demonstrated under pressure.

The Leadership Transition

On February 12, 2026, Jacob Steeves (@const_reborn) announced he was stepping down as CEO of the Opentensor Foundation (@opentensor). Co-founder Ala Shaabana (@shibshib89), known on chain as ShibShib, stepped down as COO at the same time. Steeves' own framing of the move, posted to the Bittensor Discord and quoted across community coverage, was that he would continue to write chain code, build subnets, run Novelty Search, and participate in calls, but no longer with the same legal authority. The shift was deliberate, framed as a step toward what Steeves called a "headless protocol," part of a broader push to migrate governance authority off the Foundation's balance sheet and onto the chain.

Bittensor's own governance documents describe the network as operating under a transitional Triumvirate of Opentensor Foundation employees holding root permissions alongside a senate, rather than a fully open governance model.

The Triumvirate structure has been the single most consistent criticism leveled at the network since mainnet launched in January 2021. The February resignation was the first formal step toward closing that concentration. Whether the change in legal authority would translate into a change in operating practice was the question the next 60 days would answer.

The Templar Result

On March 10, 2026, Bittensor Subnet 3, operated by the Covenant AI team under the Templar branding (@tplr_ai), announced completion of Covenant-72B. The model card on Hugging Face and the arXiv technical report confirm the specifications: a 72 billion parameter language model trained from scratch on approximately 1.1 trillion tokens of English text, using a coordination layer called Gauntlet running on Subnet 3. The training spanned more than 70 independent peer nodes, approximately 20 of which synchronized per round, each equipped with 8 B200 GPUs. The communication efficient optimizer was SparseLoCo, which compresses gradients sufficiently to make decentralized training viable across commodity internet connections rather than tightly coupled InfiniBand clusters. Weights and code were released under Apache License 2.0.

The base model scored 67.1 on MMLU zero shot, beating Meta's Llama-2- 70B at 65.6 and LLM360 K2 at 65.5. Among prior decentralized training projects, Prime Intellect's INTELLECT-1 at 100 billion parameters scored 32.7 on the same benchmark, and the whitelisted Psyche Consilience at 400 billion parameters

scored 24.2. Covenant-72B was the first decentralized training result to come within meaningful distance of a 2023-era frontier release.

The reception outside the ecosystem followed within ten days. On March 16, Anthropic (jAnthropicAI) co-founder Jack Clark (@jackclarkSF) cited the result in his research progress note under the heading

"Challenging the Political Economy of AI Through Distributed Training," explicitly flagging it as a development worth tracking even as he noted the underlying model was already two architecture generations behind frontier. On March 20, Jensen Huang of Nvidia (@nvidia) discussed the project on the All-In Podcast at the

prompt of Chamath Palihapitiya (@chamath), comparing the architecture to a modern Folding@home for AI. The Templar team's announcement post on X accumulated 1.7 million views and over 6,000 likes. SN3's token rose 194 percent in seven days. TAO itself rallied roughly 90 percent across the March cycle, moving from approximately \$170 to \$337 at peak. According to The Block's reconstruction, more than half of that price move was attributable to the Templar narrative alone.

That concentration of price action in one team's narrative was the underlying fragility the network had not yet priced.

The April 10 Exit

Sam Dare (jDistStateAndMe), founder of Covenant AI (@1Covenant), is Samuel "Sam" B. Dare, an Oxford Saïd Business School alumnus based in Dubai. His LinkedIn lists him as the founder of Covenant AI. He worked at the Opentensor Foundation alongside Steeves before transitioning to build on top of the network through Covenant. Covenant operated three of Bittensor's highest emission subnets: Templar (SN3) for pretraining, Basilica (SN39) for decentralized compute, and Grail (SN81) for reinforcement learning and post-training work.

Running one active subnet at scale requires meaningful commitment. Running three placed Covenant among the most operationally significant participants in the entire ecosystem.

On April 10, 2026, the operation ended. The sequence is documented in detail across The Block, Cointelegraph, Bitcoin.com News, Yahoo Finance, and on chain via Taostats.

Before Dare's public statement, on chain data shows approximately 37,000 TAO of subnet alpha tokens were sold from Covenant-controlled wallets across SN3, SN39, and SN81. At prevailing prices, the dump was worth approximately \$10 to \$11 million. TAO's overall trading volume in the 24 hours before the public letter hit its highest level since December 2024, which is the strongest available signal that the exit was timed for maximum extraction rather than driven by spontaneous principle. Retail holders of Covenant's three subnet tokens absorbed the impact before they had any warning.

The public letter, when it landed, framed the exit as a principled departure over governance. Dare accused Steeves of operating Bittensor under what he termed "decentralization theatre," with four specific allegations: that Steeves had suspended emissions to Covenant's subnets, removed Covenant's moderation capabilities over its own community channels, unilaterally deprecated Covenant's subnet infrastructure, and applied economic pressure through large, timed token sales. Dare's broader claim was that Bittensor's Triumvirate structure, with three signers on the upgrade multisig, was a legal shield, and that Steeves retained effective control regardless of the formal arrangement.

On the same day, a website called Tao Papers launched. Its central forensic claim was that of 41 runtime upgrades deployed between 2023 and 2026, 38 had been proposed, first-signed, and deployed from infrastructure controlled by Steeves, with the other two signers co-signing within minutes and no public discussion thread referenced. The

site also alleged that seven other subnets had emissions throttled or zeroed within 72 hours of their owners publicly disagreeing with Steeves on Discord, X, or in private Telegram groups it had been given access to. The Tao Papers data has not been publicly refuted with counter-forensics. The simultaneous launch of the site and the open letter signals that the campaign was coordinated rather than independently sourced.

The market reaction was immediate. By Friday evening, TAO had fallen from approximately \$337 to \$253 in roughly six hours, a drop of 25 percent. Market capitalization wiped out was approximately \$650 million. Long liquidations totaled \$9.1 million. Trading volume on April 10 hit \$1.72 billion, more than triple the monthly average. Subnet tokens fell harder than TAO on contagion: per BlockchainReporter's intraday tracking, Grail (SN81) was down 68 percent, Basilica (SN39) down 61.7 percent, Templar (SN3) down 57.5 percent, with Minotaur, Targon, TrajectoryRL, and NexisGen AI each falling between 26 and 34 percent. Spot outflows exceeded \$70 million across multiple consecutive days. Templar's alpha token, which had traded above \$44 at the March peak, fell into the \$5 to \$6 range and has not recovered.

What Steeves Addressed, and What He Did Not

Steeves responded the same day in a series of X posts and more fully in a follow-up Discord session over the next 72 hours. The technical rebuttal was specific and largely held up on the narrow allegations.

On emissions, Steeves stated that he has no admin authority to pause subnet emissions. Under Dynamic TAO, which launched in February 2025, emissions are determined entirely by net TAO flows into each subnet's liquidity pool through an on-chain algorithm. His own alpha sales in Covenant's subnets affected the flow based emission calculation in the same way any TAO holder's buys or sells affect it. He disclosed that he had sold less than 1 percent of what he had personally invested in Covenant's teams, and that the affected subnets at the time were running on near 100 percent burn code rather than producing active work. The Dynamic TAO emission mechanism is a documented protocol feature. The specific allegation that Steeves can unilaterally suspend emissions does not match the protocol's actual capabilities.

On moderation, Steeves acknowledged temporarily removing Dare's ability to delete posts in the Bittensor Discord and then reinstating it,

while keeping Dare's moderator role intact. Community member Alex DRocks (jDrocksAlex2) corroborated in real time that it was Dare who had deprecated his own Discord channels and that the deleted posts were community questions critical of SN39 replicating work that another compute subnet was already doing. That account was publicly available and went unchallenged by Dare in the days after.

On the personal dimension, Steeves described Dare as someone he had "considered a brother" and characterized the manner of the exit as designed to inflict "maximum pain." He apologized to holders who experienced what he called a "financial rug" and acknowledged that the network had a genuine governance vulnerability the incident had exposed. Mark Jeffrey (jMarkJeffrey) of Stillcore Capital summarized community response: "Bittensor is quite a lot more than Subnet 3, and TAO will carry on fine without it."

What Steeves did not address directly was the Tao Papers structural critique. The forensic claim that 38 of 41 network upgrades originated from infrastructure controlled by one founder, with co-signers approving within

minutes and no public discussion record, was not refuted. The earlier February resignation from the CEO role is the strongest defense available against that claim, but it does not erase the historical pattern the Tao Papers documented, and Steeves' continued role as a contributor and code author means his on-chain influence remains substantial even without the legal authority of the executive title.

That gap, between the specific allegations Steeves successfully rebutted and the structural pattern he did not address, is the honest read on the dispute.

Conviction Ships May 13

On April 16, in a Discord AMA, Steeves introduced the Conviction Mechanism, formalized shortly after as BIT-0011 under the title "Locked Stake and Conviction." The protocol ships to mainnet on May 13, 2026.

The mechanics are straightforward. From May 13 forward, every emission a subnet owner earns is auto-locked the moment it arrives in the wallet. Owners can voluntarily lock additional alpha on top of the auto-lock to build an on-chain record of commitment. The Conviction score is calculated as Stake multiplied by Time and recalculated every 30 days through an exponential moving average. Unlocking requires a publicly visible on-chain transaction. After roughly 30 days of unlock

initiation, 63 percent of the unlocked tokens become spendable. After roughly 90 days, 95 percent.

Phase one functions as an exit warning system. The moment a subnet owner initiates an unlock transaction, every staker on the network sees it within seconds and has roughly 30 days to decide whether to hold, lock their own tokens toward a competing candidate, or exit before the founder does.

Phase two converts the locked Conviction data into transferable governance power, giving the community a stake-weighted mechanism to remove underperforming owners and install replacements.

Steeves acknowledged that the locked stake mechanism was among the last pieces of work Sam Dare himself drafted before leaving the Opentensor Foundation. The failure to implement it sooner, Steeves wrote, was his "real error." The protocol fix to the Covenant crisis was drafted by the man who exploited the vulnerability the fix is designed to close.

The deeper architectural shift sits behind Conviction. Steeves' immediate response to the exit was that the crisis would "prove to birth the first subnets on Bittensor that run headless and as true commodities." Headless subnets are designed to operate without a single operator as a point of failure.

In the case of Covenant, this is exactly what happened in practice without any protocol upgrade. Community miners, with no central operator coordination, rebuilt SN3, SN39, and SN81 from open source code within days. The chain kept routing emissions throughout. New teams took on operator roles. The replacement team on what was previously Templar is now operating under the Teutonic brand and is targeting a 1 trillion parameter decentralized training run for mid to late May.

That target is the closest forward-looking milestone on the network's calendar. If Teutonic delivers, it would be by an order of magnitude the largest decentralized training run ever attempted, and it would land directly inside the August 2026 SEC review window for the spot TAO ETF filings. If Teutonic slips or the result lands below

frontier expectations, the question of whether frontier capability walked out with Covenant becomes harder to answer in the network's favor.

The Output Layer Kept Building

The ecosystem did not pause for Covenant. Several developments moved the network materially in the same window the governance crisis was unfolding.

The December 2025 halving cut daily emissions from 7,200 to 3,600 TAO. According to public on-chain trackers and Grayscale Research's pre-halving note, the event landed between December 13 and 14, with sources varying on the exact block. Inflation dropped from approximately 26 percent annualized to approximately 13 percent. The halving combined with Dynamic TAO's flow-based emission allocation has compressed supply meaningfully while approximately 68 to 70 percent of circulating supply remains locked in staking.

On April 7, Grayscale (jGrayscale) increased TAO's weighting in its AI focused fund to 43.06 percent, the largest single rebalance in the fund's history. That allocation predated the Covenant exit and persisted through it. On April 20, BitGo (jBitGo) and Yuma (jYumaGroup) launched institutional grade custody and staking for subnet alpha tokens, the first regulated rail allowing accredited capital to participate in subnet economics directly rather than only through TAO exposure. Grayscale and Bitwise both have spot TAO ETF applications pending with the SEC, with review expected around August.

On May 5, Wormhole Labs (@wormhole) activated canonical TAO bridging to Solana through its Sunrise platform, putting the asset on a second major chain and into the Jupiter Exchange flow.

On May 8, Yuma launched TaonSquare at taonsquare.com, an MCP server and AI app directory providing structured access to the full subnet catalog: what each subnet offers, pricing, API availability, and market cap. Yuma is the DCG-backed (jDCGco) Bittensor accelerator co-led by Barry Silbert (jBarrySilbert) and now Greg Schvey (jGSchvey) and Jeff Schvey, and is also behind Yuma Asset Management's Subnet Composite and Large Cap funds. The accessibility layer the magazine has been writing about for months is finally receiving serious attention, and it is being built by an accelerator rather than by the Opentensor Foundation. That distinction matters when reading who is doing the work of making the network legible to capital.

On May 9, Grayscale reopened private placements for the GDAO Trust to accredited investors. On May 10, Subnet 68, operated by Metanova Labs (@metanova_labs), hit a milestone of screening 11 million drug molecules across nine disease targets in live competitions. Metanova also

announced a partnership with ONEPOT.AI to move AI-screened compounds to robotic synthesis in days rather than weeks. Manako is building AI vision agents on Score (SN44) for enterprise security camera

operational intelligence. Apex (SN1), the flagship language model subnet with agentic workflows and tool use, launched RL Tron, a reinforcement learning competition where AI agents face off in head-to-head Tron tournaments. Targon launched the Targon Supply Portal allowing compute suppliers to monetize idle hardware. NIOME (SN55) won the MIT Entrepreneurial Development Prize in January and secured a commercial client in Dr. Monika Gostic Nutrition. Astrid Intelligence published an analysis arguing that the

Covenant volatility did not disrupt core network operations.

Chutes (@chutes_ai, SN64), the serverless inference subnet, remains the closest drop-in replacement for the OpenAI API that exists in crypto AI today. Hippus (SN75) is offering distributed cloud storage at a unit cost roughly 2,386 times below the major hyperscalers. Ridges (@ridges_ai, SN62) is performing competitively against Cursor on autonomous software engineering benchmarks while trading at approximately \$30 million in market capitalization against Cursor's \$30 billion private valuation.

The output layer kept building through the worst governance crisis the network has experienced.

Where the Market Sits

As of May 11, 2026, TAO traded at approximately \$317.86 according to CoinDesk, with other aggregators reporting between \$317 and \$327 depending on venue. Market capitalization sits at approximately \$3 billion to \$3.5 billion. The asset is roughly 57 percent below its April 2024 all time high of \$760. The 24 hour trading range on May 11 spanned \$309 to \$333. Q1 2026 network revenue was \$43 million.

The broader crypto market context matters. Bitcoin opened May 11 at \$82,164, the strongest opening since January 31, before sliding back to roughly \$81,000 by midday. Ethereum traded around \$2,330 to \$2,370. Bitcoin had crashed 30 percent in late January and has been consolidating ever since, weighed down by Middle East geopolitical risk, \$400 million plus in spot BTC ETF outflows in late April, and reduced near-term rate cut expectations after the Fed held at 3.50 to 3.75 percent. TAO's recovery from \$250 back to \$317 has happened against a choppy broader crypto tape, which is at least a partial

indication that the recovery is being driven by network-specific catalysts rather than beta to the market.

The Centralized Side

Bittensor's governance concentration was exposed under hostile pressure. The closed lab economy has a different and equally serious accounting problem, and reading one without the other produces a misleading picture of where intelligence production is actually paying for itself.

According to Alphabet's own Q1 2026 10-Q filing with the SEC, the company reported \$36.9 billion in pre-tax gains on equity securities. The net contribution to net income after tax was \$28.7 billion, adding \$2.35 to diluted earnings per share. Alphabet's total Q1 net income of \$62.6 billion is therefore roughly 46 percent attributable to the markup of its private equity stakes, primarily its 13 to 14 percent stake in Anthropic. Amazon disclosed \$16.8 billion in pre-tax gains from its own Anthropic investment in the same quarter, more than half of Amazon's pre-tax income. Both companies are required by GAAP to mark their stakes to the most recent funding round price. Anthropic itself is reportedly reviewing offers between \$800 and \$900 billion just 76 days after closing a \$380 billion Series G in February.

When two of the largest cloud providers in the world earn roughly half of their reported AI profits from marking up an equity stake in a private company they themselves are committed to spending tens of billions of dollars with over the next five years, the relationship between accounting income and operating income is not what the headline numbers suggest. The Anthropic \$200 billion Google commitment, signed May 5, formalizes the circular flow: Google invests up to \$40 billion in Anthropic, Anthropic commits \$200 billion back to Google for

TPU compute and cloud services beginning in 2027 across 5 gigawatts of capacity, and Alphabet books the implied valuation increase as quarterly profit. According to The Information's reporting, Anthropic and OpenAI together now account for more than half of the \$2 trillion in long term cloud backlog held across AWS, Azure, GCP, and Oracle.

The pure-play labs' own loss trajectories make the picture clearer. Anthropic moved from approximately \$9 billion to approximately \$30 billion in annualized run rate between January and April 2026, with internal documents projecting a \$14 billion operating loss in 2026 and no positive free cash flow until 2028 at earliest. OpenAI is at roughly \$24

billion ARR at a reported \$852 billion valuation. Leaked internal projections reported by the Wall Street Journal show a \$74 billion operating loss in 2028 alone. The growth is real. The unit economics have not budged.

Meanwhile, the open weight cohort compresses the closed lab moat from below. DeepSeek V4-Pro sits at 80.6 percent on SWE Bench Verified, within 0.2 points of Claude Opus 4.6, released under MIT license at approximately one tenth to one thirteenth the API cost. Qwen 3.6, Llama 4, GLM 5.1, Kimi K2.6, and Gemma 4 each reached frontier parity within months of corresponding closed releases. The ceiling on what frontier APIs can charge is set by what it costs the customer to self-host an open weight equivalent, and that ceiling is dropping 30 to 50 percent a year.

This is the structural problem Bittensor's thesis was always pointing at, and it has not gone away because of the Covenant exit. The closed labs are running a system whose unit economics depend on circular capital flows between hyperscalers and the labs they invest in, an open weight cohort that compresses the API ceiling from below, and revenue growth that has not yet caught up to the cost base. Decentralized AI does not have these problems because there is no equity stake to mark up and no balance sheet on which to take a \$74 billion operating loss. There is only the token, and the token reflects what the network actually produces.

The question is whether the network produces enough.

Reading the Network Through May 12

Bittensor (@opentensor) in May 2026 is a network with measurable revenue at \$43 million for Q1, real builders shipping product across pretraining, inference, compute, storage, fine-tuning, drug discovery, vision agents, and reinforcement learning, an institutional rail layer that survived the worst governance crisis in the network's history without significant net outflows, a community that rebuilt three of the most operationally important subnets from open source code in days, a protocol response in BIT-0011 that ships to mainnet on May 13, and a frontier capability question that Teutonic's planned 1 trillion parameter run will either answer or expose by late May.

It is also a network whose former CEO and his co-founder formally stepped down from executive roles in February precisely because the governance concentration the Tao Papers documented was a real

problem, whose accessibility layer is being built by a DCG-backed accelerator rather than by the protocol, whose subnet alpha token aggregate market cap sits at approximately \$1.03 billion or roughly 30 percent of

TAO's own valuation, and whose price chart still reflects the rotation-driven nature of crypto AI rather than the operating leverage of an intelligence marketplace. Three things have to hold in mind simultaneously to read the network honestly. The Templar/Covenant-72B result was a real proof point that decentralized substrate can produce frontier scale training runs, and the team that produced it is no longer on the network and exited in a manner that has rightly damaged community trust in star-team valuation frameworks. The Covenant exit was the most damaging single event in the network's history and it also accelerated the protocol response that was already in development. The ecosystem around the crisis kept building, and the institutional infrastructure rails kept extending.

The closed labs are not solving for the same constraints. Alphabet booked \$28.7 billion of last quarter's profit by marking up its Anthropic stake. Anthropic signed a \$200 billion commitment to spend that money back. The system functions because capital keeps flowing into a circle that has not yet been required to produce positive unit economics on its own. The decentralized model has no such backstop. It does not have an equity stake to mark up. It does not have a Series G to draw down. It has the token, the subnet incentives, and whatever output the network produces. That is a harder problem on a smaller scale. It is also a problem whose solution, if it arrives, compounds in a way the closed system's solution cannot.

May 13 is tomorrow and Conviction ships. Late May is Teutonic and the trillion parameter target. August is the SEC ETF review window. The data will arrive on a calendar, and the question for the next four quarters is whether Bittensor's structural advantage over the closed lab economy compounds faster than the structural disadvantage of a governance layer that has only just begun to mature.

Ninety days ago, the network had a star team and a thesis. Today it has a stress-tested incentive layer, a community that proved it can route around operator failure, a protocol fix shipping in 24 hours, and a frontier run on the calendar. The thesis is not yet validated. It is no longer untested either.

10:28 PM · May 12, 2026 · 101 Views

Post your reply Reply